

Characterising effect of anaemia on mortality in severe malaria

Contents

Background	1
Imputation of missing variables	1
Exploratory analysis	3
Predictive value of anaemia on death adjusting for confounders	8
More complex GAM model	10
Model comparison	12

Background

This looks at the severe malaria legacy dataset from MORU

Imputation of missing variables

Quite a lot of the important covariates are missing in the older studies. We use linear regression to estimate these unknown variables:

- Missing base deficit is imputed using bicarbonate (if available) else using respiratory rate
- Missing Blood urea nitrogen is imputed using creatinine

Impute base deficit from bicarbonate

```
BD_and_bicarbonate = !is.na(Leg_data$BD) & !is.na(Leg_data$bicarbonate)
print(paste('We have ', sum(BD_and_bicarbonate), 'observations for both bicarbonate and base deficit'))

## [1] "We have 5067 observations for both bicarbonate and base deficit"

mod_impute1 = lmer(BD ~ bicarbonate + (1 | studyID) + (1 | country), data= Leg_data[BD_and_bicarbonate,])
missing_BD = is.na(Leg_data$BD)
Available_Bicarbonate = !is.na(Leg_data$bicarbonate)
print(paste(sum(missing_BD & Available_Bicarbonate), 'observations will now be imputed'))

## [1] "309 observations will now be imputed"

# impute with model
Leg_data$BD[missing_BD & Available_Bicarbonate] = predict(mod_impute1,newdata=Leg_data[missing_BD & Available_Bicarbonate,])
```

Impute base deficit from lactate

```
BD_and_lactate = !is.na(Leg_data$BD) & !is.na(Leg_data$lactate)
print(paste('We have ', sum(BD_and_lactate), 'observations for both lactate and base deficit'))

## [1] "We have 632 observations for both lactate and base deficit"
```

```

if(length(unique(Leg_data$studyID[BD_and_lactate]))==1){
  mod_impute2 = lm(BD ~ lactate, data= Leg_data[BD_and_lactate,])
} else {
  mod_impute2 = lmer(BD ~ lactate + (1 | studyID), data= Leg_data[BD_and_lactate,])
}
missing_BD = is.na(Leg_data$BD)
Available_Lactate = !is.na(Leg_data$lactate)
print(paste(sum(missing_BD & Available_Lactate), 'observations will now be imputed'))

```

```
## [1] "722 observations will now be imputed"
```

```
# impute with model
```

```
Leg_data$BD[missing_BD & Available_Lactate] = predict(mod_impute2,newdata=Leg_data[missing_BD & Available_Lactate,])
```

Impute base deficit from respiratory rate

```

BD_and_rr = !is.na(Leg_data$BD) & !is.na(Leg_data$rr)
print(paste('We have ', sum(BD_and_rr), 'observations for both resp rate and base deficit'))

```

```
## [1] "We have 7572 observations for both resp rate and base deficit"
```

```

mod_impute3 = lmer(BD ~ rr + (1 | studyID), data= Leg_data[BD_and_rr,])
missing_BD = is.na(Leg_data$BD)
Available_rr = !is.na(Leg_data$rr)
print(paste(sum(missing_BD & Available_rr), 'observations will now be imputed'))

```

```
## [1] "1650 observations will now be imputed"
```

```
Leg_data$BD[missing_BD & Available_rr] = predict(mod_impute3,newdata=Leg_data[missing_BD & Available_rr,])
```

Impute blood urea nitrogen from creatinine:

```

BUN_and_cr = !is.na(Leg_data$BUN) & !is.na(Leg_data$creatinine)
print(paste('We have ', sum(BUN_and_cr), 'observations for both blood urea nitrogen and creatinine'))

```

```
## [1] "We have 1453 observations for both blood urea nitrogen and creatinine"
```

```

mod_impute4 = lmer(BUN ~ creatinine + (1 | studyID), data= Leg_data[BUN_and_cr,])
missing_BUN = is.na(Leg_data$BUN)
Available_cr = !is.na(Leg_data$creatinine)
print(paste(sum(missing_BUN & Available_cr), 'observations will now be imputed'))

```

```
## [1] "679 observations will now be imputed"
```

```
Leg_data$BUN[missing_BUN & Available_cr] = predict(mod_impute4,newdata=Leg_data[missing_BUN & Available_cr,])
```

Resulting data we can now use: The contributions of the different studies:

```

vars_interest = c('outcome','HCT','LPAR_pct','BD','BUN','poedema',
                  'convulsions','coma','AgeInYear','drug_class')
complete_cases = apply(Leg_data[,vars_interest], 1, function(x) sum(is.na(x))) == 0
Complete_Leg_data = Leg_data[complete_cases,] # for the model fitting
Complete_Leg_data$studyID = as.factor(as.character(Complete_Leg_data$studyID))
# Whole dataset
table(Leg_data$studyID)

```

```
##
```

```
##          AAV          AQ      AQGambia      AQUAMAT Core Malaria
##          370          560          579          5494          1122
```

```
## SEAQUAMAT
```

```
##          1461
# in the complete dataset (all variables recorded)
table(Complete_Leg_data$studyID)

##
##          AAV          AQ          AQGambia          AQUAMAT Core Malaria
##          214          150          168          3666          657
##    SEQUAMAT
##          1333

Complete_Leg_data$drug_AS = 0
Complete_Leg_data$drug_AS[Complete_Leg_data$drug_class=='artemisinin']=1

# remove infinite log parasitaemias
ind_keep = !(is.infinite(Complete_Leg_data$LPAR_pct) | is.nan(Complete_Leg_data$LPAR_pct))
Complete_Leg_data = Complete_Leg_data[ind_keep,]
```

Exploratory analysis

```
for(s in unique(Complete_Leg_data$studyID)){
  print(paste(s, ', mortality of:', round(100*mean(Complete_Leg_data$outcome[Complete_Leg_data$studyID==s]), 1)))
}

## [1] "Core Malaria , mortality of: 23 %"
## [1] "AQGambia , mortality of: 12 %"
## [1] "AAV , mortality of: 12 %"
## [1] "SEQUAMAT , mortality of: 18 %"
## [1] "AQUAMAT , mortality of: 9 %"
## [1] "AQ , mortality of: 23 %"

for(s in unique(Complete_Leg_data$studyID)){
  print(paste0(s, ', ages:', round(quantile(Complete_Leg_data$AgeInYear[Complete_Leg_data$studyID==s]), 1)))
}

## [1] "Core Malaria, ages:1Core Malaria, ages:27Core Malaria, ages:75"
## [1] "AQGambia, ages:1AQGambia, ages:4AQGambia, ages:9"
## [1] "AAV, ages:15AAV, ages:34AAV, ages:77"
## [1] "SEQUAMAT, ages:2SEQUAMAT, ages:25SEQUAMAT, ages:87"
## [1] "AQUAMAT, ages:0AQUAMAT, ages:2AQUAMAT, ages:78"
## [1] "AQ, ages:15AQ, ages:30AQ, ages:74"

for(s in unique(Complete_Leg_data$studyID)){
  print(s)
  print(table(Complete_Leg_data$drug[Complete_Leg_data$studyID==s]))
}

## [1] "Core Malaria"
##
##    Amodiaquine    Artemether    Artesunate    Chloroquine    Lumefantrine
##           0           11           368           2           0
##    Mefloquine      NAC      Quinine
##           7           6           262
## [1] "AQGambia"
##
```

```

## Amodiaquine Artemether Artesunate Chloroquine Lumefantrine
##      0      82      0      0      0
## Mefloquine      NAC      Quinine
##      0      0      86
## [1] "AAV"
##
## Amodiaquine Artemether Artesunate Chloroquine Lumefantrine
##      0      102      112      0      0
## Mefloquine      NAC      Quinine
##      0      0      0
## [1] "SEAQUAMAT"
##
## Amodiaquine Artemether Artesunate Chloroquine Lumefantrine
##      0      0      645      0      0
## Mefloquine      NAC      Quinine
##      0      0      628
## [1] "AQUAMAT"
##
## Amodiaquine Artemether Artesunate Chloroquine Lumefantrine
##      0      0      1837      0      0
## Mefloquine      NAC      Quinine
##      0      0      1818
## [1] "AQ"
##
## Amodiaquine Artemether Artesunate Chloroquine Lumefantrine
##      0      73      0      0      0
## Mefloquine      NAC      Quinine
##      0      0      77

```

Let's look at the key predictive variables. We use a random effects term to model differences between studies.

```

## Linear mixed model fit by REML ['lmerMod']
## Formula: BD ~ HCT + (1 | studyID/country)
## Data: Complete_Leg_data
##
## REML criterion at convergence: 40261.9
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -4.4421 -0.6612 -0.1488  0.5224  4.7209
##
## Random effects:
## Groups      Name      Variance Std.Dev.
## country:studyID (Intercept) 2.6525  1.6286
## studyID      (Intercept) 0.8373  0.9151
## Residual                    41.8947  6.4726
## Number of obs: 6116, groups: country:studyID, 18; studyID, 6
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept) 10.339058  0.653393  15.82
## HCT         -0.133548  0.009699 -13.77
##
## Correlation of Fixed Effects:
##      (Intr)

```

```

## HCT -0.394

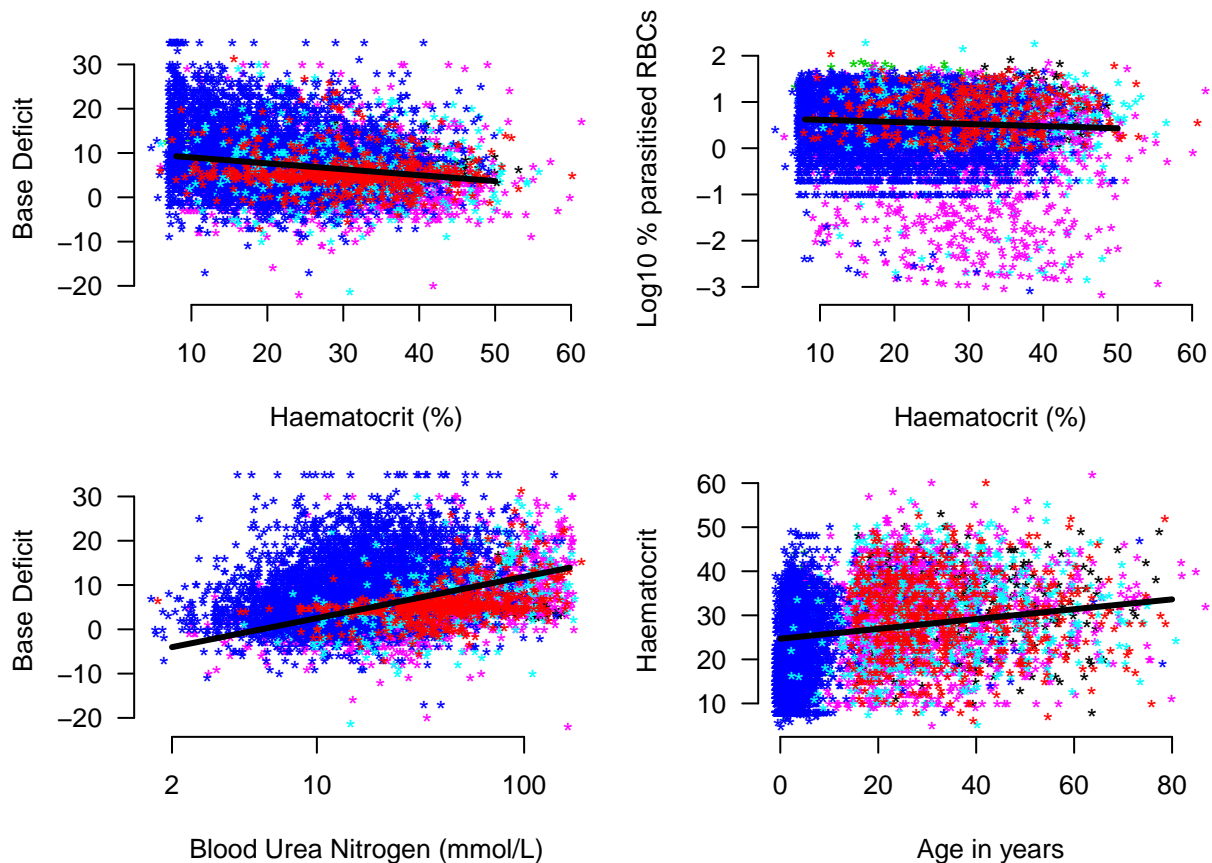
## Linear mixed model fit by REML ['lmerMod']
## Formula: LPAR_pct ~ HCT + (1 | studyID/country)
## Data: Complete_Leg_data
##
## REML criterion at convergence: 13822.9
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -4.7144 -0.5555  0.1598  0.7265  2.4355
##
## Random effects:
##      Groups             Name             Variance Std.Dev.
## country:studyID (Intercept)  0.00946   0.09726
## studyID          (Intercept)  0.07496   0.27379
## Residual                                0.55564   0.74541
## Number of obs: 6116, groups:  country:studyID, 18; studyID, 6
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)  0.659944   0.121244   5.443
## HCT          -0.004579   0.001116  -4.105
##
## Correlation of Fixed Effects:
##      (Intr)
## HCT -0.251

## Linear mixed model fit by REML ['lmerMod']
## Formula: BD ~ log10(BUN) + (1 | studyID/country)
## Data: Complete_Leg_data
##
## REML criterion at convergence: 39236.2
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -5.6063 -0.6369 -0.1041  0.5191  5.0754
##
## Random effects:
##      Groups             Name             Variance Std.Dev.
## country:studyID (Intercept)  2.876   1.696
## studyID          (Intercept)  6.858   2.619
## Residual                                35.405   5.950
## Number of obs: 6116, groups:  country:studyID, 18; studyID, 6
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)  -6.8409   1.2574  -5.44
## log10(BUN)    9.3530   0.2559  36.55
##
## Correlation of Fixed Effects:
##      (Intr)
## log10(BUN) -0.293

## Linear mixed model fit by REML ['lmerMod']

```

```
## Formula: HCT ~ AgeInYear + (1 | studyID/country)
## Data: Complete_Leg_data
##
## REML criterion at convergence: 43534.9
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.1004 -0.7399 -0.0515  0.6927  3.5627
##
## Random effects:
## Groups      Name      Variance Std.Dev.
## country:studyID (Intercept)  5.722   2.392
## studyID      (Intercept)  7.322   2.706
## Residual                71.467   8.454
## Number of obs: 6116, groups: country:studyID, 18; studyID, 6
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept) 24.69246   1.36141  18.137
## AgeInYear    0.11159   0.01159   9.626
##
## Correlation of Fixed Effects:
##              (Intr)
## AgeInYear -0.185
```



If we look at the relationship between haematocrit and death:


```

par(las=1, bty='n')
Complete_Leg_data$country=as.factor(Complete_Leg_data$country)
modHCT=gam(outcome ~ s(HCT) + s(studyID, bs='re') + s(country, bs='re'),data = Complete_Leg_data, family=
summary(modHCT)

##
## Family: binomial
## Link function: logit
##
## Formula:
## outcome ~ s(HCT) + s(studyID, bs = "re") + s(country, bs = "re")
##
## Parametric coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.8865      0.2397   -7.87 3.54e-15 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##             edf Ref.df  Chi.sq p-value
## s(HCT)        2.304  2.922   5.482 0.15478
## s(studyID)     3.611  5.000 314.182 0.00361 **
## s(country)    10.766 14.000 162.027 0.01464 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.0484   Deviance explained = 5.99%
## UBRE = -0.25846   Scale est. = 1           n = 6116

preds = predict(modHCT, newdata = data.frame(HCT=5:45, studyID='AQ', country='Thailand', country=1),
                exclude = c("s(country)", "s(studyID)"), type='response', se.fit=T)
plot(5:45, 100*preds$fit, ylab='Probability death (%)', xlab='Haematocrit',
     type='l', lwd=3, ylim = c(5,25))
lines(5:45, 100*preds$fit + 100*2*preds$se.fit, lty=2, lwd=2)
lines(5:45, 100*preds$fit - 100*2*preds$se.fit, lty=2, lwd=2)

```



Predictive value of anaemia on death adjusting for confounders

Before fitting the more complex GAM models we explore the standard glm (logistic regression) models.

```
mod_full_GLM = glmer(outcome ~ HCT + LPAR_pct + AgeInYear + coma + convulsions +
  poedema + log10(BUN) + BD + drug_AS +
  (1 | studyID) + (1 | country),
  data = Complete_Leg_data, family=binomial)
```

```
## Warning in checkConv(attr(opt, "derivs"), opt$par, ctrl = control
## $checkConv, : Model failed to converge with max|grad| = 0.00101495 (tol =
## 0.001, component 1)
```

```
summary(mod_full_GLM)
```

```
## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula:
## outcome ~ HCT + LPAR_pct + AgeInYear + coma + convulsions + poedema +
## log10(BUN) + BD + drug_AS + (1 | studyID) + (1 | country)
## Data: Complete_Leg_data
##
##      AIC      BIC  logLik deviance df.resid
## 3459.7  3540.3 -1717.8  3435.7     6104
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.9034 -0.3324 -0.1914 -0.1076 15.4072
##
## Random effects:
## Groups  Name              Variance Std.Dev.
```



```
## country (Intercept) 1.501e-01 3.875e-01
## studyID (Intercept) 1.919e-09 4.381e-05
## Number of obs: 6116, groups: country, 15; studyID, 6
##
## Fixed effects:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -7.000057   0.306929 -22.807 < 2e-16 ***
## HCT          0.016441   0.005284   3.111 0.001863 **
## LPAR_pct     -0.001281   0.060471  -0.021 0.983095
## AgeInYear    0.013715   0.003840   3.571 0.000355 ***
## coma        1.338046   0.100906  13.260 < 2e-16 ***
## convulsions1 0.513532   0.116864   4.394 1.11e-05 ***
## poedema1     0.543720   0.385373   1.411 0.158276
## log10(BUN)   1.778368   0.166012  10.712 < 2e-16 ***
## BD           0.121719   0.007183  16.944 < 2e-16 ***
## drug_AS      -0.343604   0.090337  -3.804 0.000143 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) HCT      LPAR_p AgInYr coma   cnvls1 poedm1 l10(BU BD
## HCT          -0.483
## LPAR_pct     -0.041  0.030
## AgeInYear    0.039 -0.178  0.001
## coma        -0.163 -0.027  0.075  0.000
## convulsins1 -0.133 -0.073  0.017  0.108 -0.220
## poedema1     -0.004 -0.005 -0.006 -0.048  0.027  0.000
## log10(BUN)   -0.700  0.064 -0.047 -0.245 -0.014  0.103  0.006
## BD           -0.148  0.199 -0.180  0.135 -0.024  0.024 -0.008 -0.262
## drug_AS      -0.091 -0.012 -0.024 -0.022  0.007  0.004 -0.025 -0.044 -0.020
## convergence code: 0
## Model failed to converge with max|grad| = 0.00101495 (tol = 0.001, component 1)
```

Now let's make counterfactual predictions of anaemia on death for the patients in the database.

```
myquantiles = c(0.25,0.5,0.75) # this is 50% predictive interval

overall_median_mortality = median(100*predict(mod_full_GLM, type='response'))
par(las=1, bty='n')
x_hcts = seq(4,45, by=1)
probs_lin = array(dim = c(3, length(x_hcts)))
for(i in 1:length(x_hcts)){
  mydata = Complete_Leg_data
  mydata$HCT=x_hcts[i]
  ys = 100*predict(mod_full_GLM, newdata = mydata, re.form=NA, type='response')
  probs_lin[,i] = quantile(ys, probs=myquantiles)
}
```


The way to interpret this 'counterfactual' plot is as follows: suppose that every individual in the dataset was assigned (as in an intervention) a specific haematocrit X , what would the resulting per patient probability of death be. Here we summarise these probabilities by the predicted mean probability of death and 80% predictive intervals.

```
plot(x_hcts,probs_lin[2,], xlim=c(4,45), ylab='Predicted probability of death',
     xlab='Haematocrit (%)', ylim=c(0,20), lty=1, lwd=3, type='l')
lines(x_hcts, probs_lin[1,], lty=2, lwd=2)
```

```

lines(x_hcts, probs_lin[3,], lty=2, lwd=2)
abline(h=overall_median_mortality, lwd=3, col='blue', lty=2)
legend('topleft', col=c('black','black','blue'), lwd=3, lty=c(1,2,2),
      legend = c('Median counterfactual value', '25th and 75th counterfactual quantiles', 'Retrodicted median value'))

```



More complex GAM model

The GAM model allows for non-linear relationships between certain variables and the outcome.

Here we fit as non-linear the effect of age and haematocrit on mortality. We add a random effect term for the studyID. We should also be doing this for the study site...

```

mod_full_GAM = gam(outcome ~ s(HCT, AgeInYear) + LPAR_pct + coma + convulsions +
                    poedema + log10(BUN) + BD + drug_AS +
                    s(studyID, bs='re') + s(country, bs='re'),
                    data=Complete_Leg_data, family=binomial)
summary(mod_full_GAM)

```

```

##
## Family: binomial
## Link function: logit
##
## Formula:
## outcome ~ s(HCT, AgeInYear) + LPAR_pct + coma + convulsions +
##           poedema + log10(BUN) + BD + drug_AS + s(studyID, bs = "re") +
##           s(country, bs = "re")
##
## Parametric coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -6.327660   0.270379 -23.403  < 2e-16 ***
## LPAR_pct      0.001763   0.060455  0.029  0.976734
## coma         1.330823   0.100889  13.191  < 2e-16 ***

```


```
## convulsions1  0.534702    0.117347    4.557  5.2e-06 ***
## poedema1     0.547874    0.384141    1.426  0.153801
## log10(BUN)   1.701427    0.170554    9.976  < 2e-16 ***
## BD          0.123330    0.007331   16.824  < 2e-16 ***
## drug_AS     -0.343908    0.090360   -3.806  0.000141 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df Chi.sq  p-value
## s(HCT,AgeInYear) 5.486902   7.664 33.629 3.61e-05 ***
## s(studyID)        0.004287   5.000  0.003   0.495
## s(country)        9.944672  14.000 74.600 2.80e-15 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.27   Deviance explained = 29.1%
## UBRE = -0.43785   Scale est. = 1          n = 6116
```

Now we compute the corresponding counterfactual probabilities of death for the dataset for all values of the haematocrit:

```
overall_median_mortalityGAM = median(100*predict(mod_full_GAM, type='response'))
par(las=1, bty='n')
probs_gam = array(dim = c(3, length(x_hcts)))
for(i in 1:length(x_hcts)){
  mydata = Complete_Leg_data
  mydata$HCT=x_hcts[i]
  ys = 100*predict(mod_full_GAM, newdata = mydata, type='response')
  probs_gam[,i] = quantile(ys, probs=myquantiles)
}
```

We see that the effect of haematocrit on mortality is non-linear under this model: below 20 is protective, above 20 plateaus out:

```
#
par(las=1, mfrow=c(1,2), bty='n', mar=c(4,4,1,1))
#### Plot the standard logistic regression model
plot(x_hcts,probs_lin[2,], xlim=c(4,45), ylab='Predicted probability of death',
     xlab='Haematocrit (%)', ylim=c(0,20), lty=1, lwd=3, type='l')
lines(x_hcts, probs_lin[1,], lty=2, lwd=2)
lines(x_hcts, probs_lin[3,], lty=2, lwd=2)
abline(h=overall_median_mortality, lwd=3, col='blue',lty=2)
title('Logistic regression model')
#### And now the GAM model
plot(x_hcts,probs_gam[2,], xlim=c(4,45), ylab='Predicted probability of death',
     xlab='Haematocrit (%)', ylim=c(0,20), lty=1, lwd=3, type='l')
lines(x_hcts, probs_gam[1,], lty=2, lwd=2)
lines(x_hcts, probs_gam[3,], lty=2, lwd=2)
abline(h=overall_median_mortalityGAM, lwd=3, col='blue',lty=2)
title('Generalised additive model')
```



Model comparison

Which model is better fit in terms of AIC

```
print(AIC(mod_full_GAM, mod_full_GLM))
```

```
##           df      AIC
## mod_full_GAM 23.43586 3438.101
## mod_full_GLM 12.00000 3459.687
```

And in terms of deviance

```
print(list(mod_full_GLM = deviance(mod_full_GLM), mod_full_GAM=deviance(mod_full_GAM)))
```

```
## $mod_full_GLM
## [1] 3400.304
##
## $mod_full_GAM
## [1] 3391.229
```